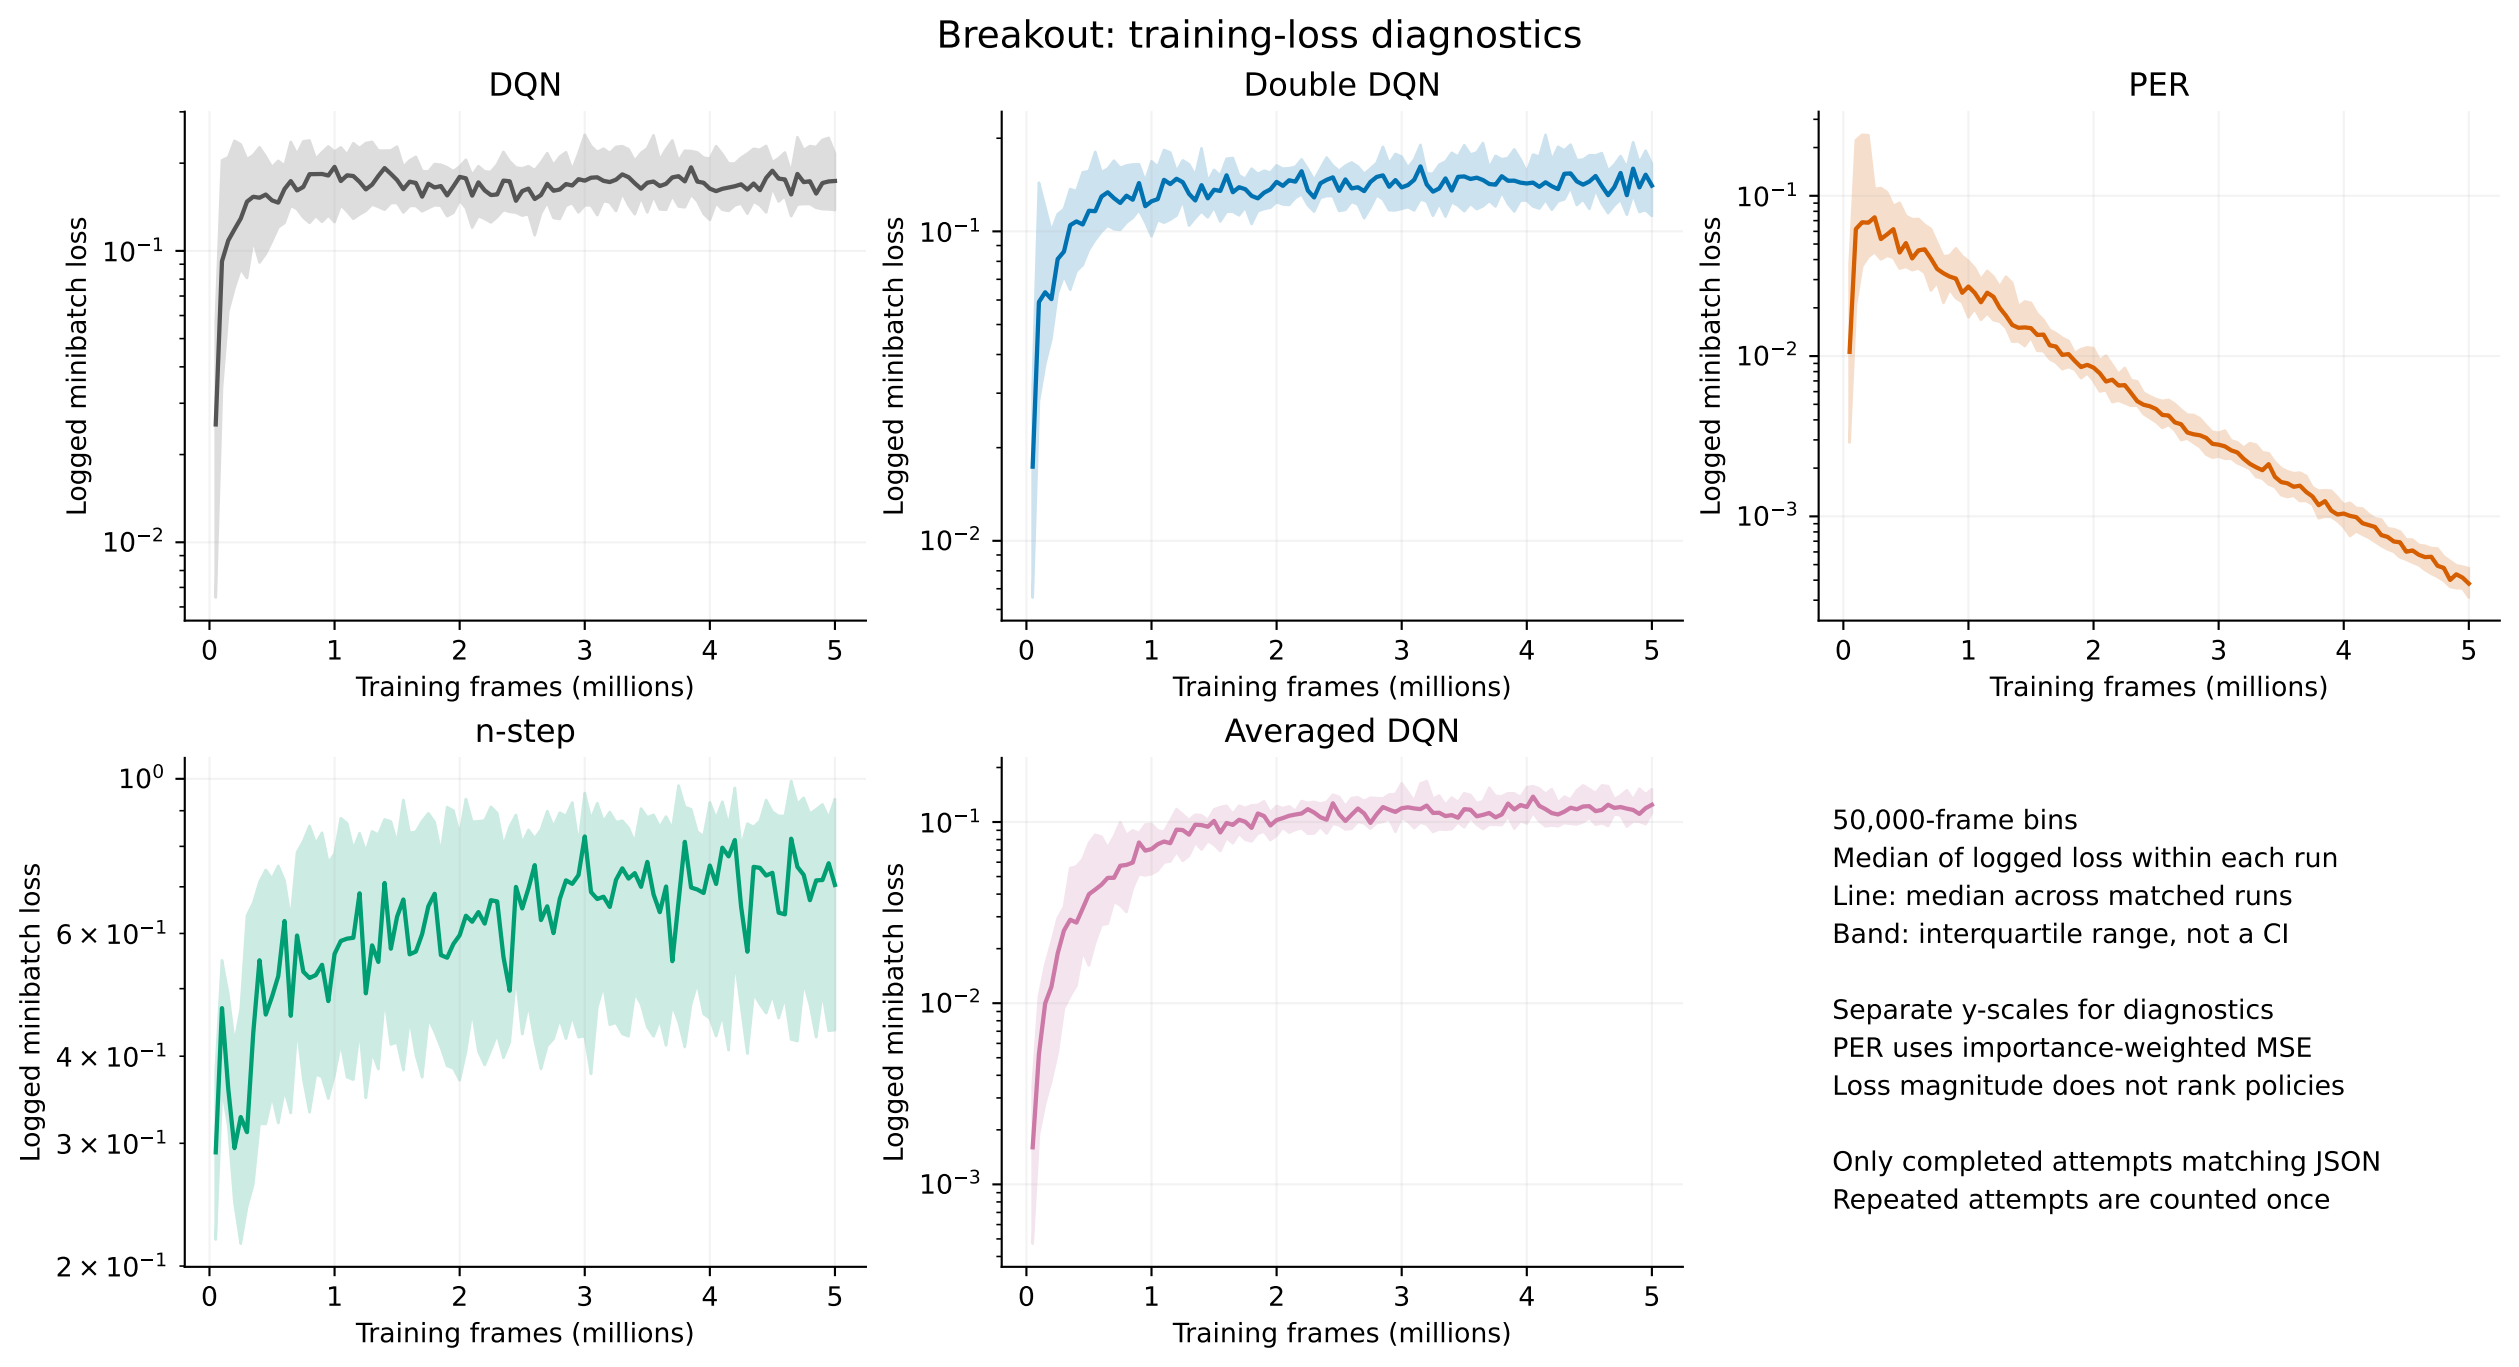
Breakout: training-loss diagnostics
DQN
Double DQN
PER
n-step
Averaged DQN
Logged minibatch loss
Training frames (millions)
50,000-frame bins
Median of logged loss within each run
Line: median across matched runs
Band: interquartile range, not a CI
Separate y-scales for diagnostics
PER uses importance-weighted MSE
Loss magnitude does not rank policies
Only completed attempts matching JSON
Repeated attempts are counted once